

dplyr tutorial

What is dplyr?

dplyr is a powerful R-package to transform and summarize tabular data with rows and columns. For another explanation of dplyr see the dplyr package vignette: Introduction to dplyr (<http://cran.rstudio.com/web/packages/dplyr/vignettes/introduction.html>)

Why is it useful?

The package contains a set of functions (or “verbs”) that perform common data manipulation operations such as filtering for rows, selecting specific columns, re-ordering rows, adding new columns and summarizing data.

In addition, dplyr contains a useful function to perform another common task which is the “split-apply-combine” concept. We will discuss that in a little bit.

How does it compare to using base functions R?

If you are familiar with R, you are probably familiar with base R functions such as `split()`, `subset()`, `apply()`, `sapply()`, `lapply()`, `tapply()` and `aggregate()`. Compared to base functions in R, the functions in dplyr are easier to work with, are more consistent in the syntax and are targeted for data analysis around data frames instead of just vectors.

How do I get dplyr?

To install dplyr

```
install.packages("dplyr")
```

To load dplyr

```
library(dplyr)
```

Data: mammals sleep

The msleep (mammals sleep) data set contains the sleeptimes and weights for a set of mammals and is available in the dagdata repository on github. This data set contains 83 rows and 11 variables.

Download the msleep data set in CSV format from here (https://raw.githubusercontent.com/genomicsclass/dagdata/master/inst/extdata/msleep_ggplot2.csv), and then load into R:

```
library(downloader)
url <- "https://raw.githubusercontent.com/genomicsclass/dagdata/master/inst/extdata/msleep_ggplot2.csv"
filename <- "msleep_ggplot2.csv"
if (!file.exists(filename)) download(url,filename)
msleep <- read.csv("msleep_ggplot2.csv")
head(msleep)
```

```
##           name      genus vore      order conservation
## 1      Cheetah  Acinonyx carni   Carnivora          lc
## 2      Owl monkey    Aotus  omni   Primates         <NA>
## 3  Mountain beaver Aplodontia herbi   Rodentia          nt
## 4 Greater short-tailed shrew Blarina  omni Soricomorpha          lc
## 5              Cow        Bos herbi Artiodactyla domesticated
## 6  Three-toed sloth  Bradypus herbi   Pilosa         <NA>
##  sleep_total sleep_rem sleep_cycle awake brainwt  bodywt
## 1      12.1      NA      NA  11.9      NA  50.000
## 2      17.0      1.8      NA   7.0 0.01550   0.480
## 3      14.4      2.4      NA   9.6      NA   1.350
## 4      14.9      2.3  0.1333333   9.1 0.00029   0.019
## 5       4.0      0.7  0.6666667  20.0 0.42300 600.000
## 6      14.4      2.2  0.7666667   9.6      NA   3.850
```

The columns (in order) correspond to the following:

column name	Description
name	common name
genus	taxonomic rank
vore	carnivore, omnivore or herbivore?
order	taxonomic rank
conservation	the conservation status of the mammal
sleep_total	total amount of sleep, in hours
sleep_rem	rem sleep, in hours
sleep_cycle	length of sleep cycle, in hours
awake	amount of time spent awake, in hours
brainwt	brain weight in kilograms
bodywt	body weight in kilograms

Important dplyr verbs to remember

dplyr verbs	Description
<code>select()</code>	select columns
<code>filter()</code>	filter rows
<code>arrange()</code>	re-order or arrange rows
<code>mutate()</code>	create new columns
<code>summarise()</code>	summarise values
<code>group_by()</code>	allows for group operations in the “split-apply-combine” concept

dplyr verbs in action

The two most basic functions are `select()` and `filter()` which selects columns and filters rows, respectively.

Selecting columns using `select()`

Select a set of columns: the name and the sleep_total columns.

```
sleepData <- select(msleep, name, sleep_total)
head(sleepData)
```

```
##              name sleep_total
## 1          Cheetah      12.1
## 2      Owl monkey      17.0
## 3  Mountain beaver      14.4
## 4 Greater short-tailed shrew      14.9
## 5              Cow       4.0
## 6  Three-toed sloth      14.4
```

To select all the columns *except* a specific column, use the “-” (subtraction) operator (also known as negative indexing)

```
head(select(msleep, -name))
```

```
##      genus vore      order conservation sleep_total sleep_rem
## 1  Acinonyx carni   Carnivora          lc        12.1        NA
## 2    Aotus  omni   Primates        <NA>        17.0         1.8
## 3 Aplodontia herbi   Rodentia          nt        14.4         2.4
## 4   Blarina  omni Soricomorpha          lc        14.9         2.3
## 5      Bos herbi Artiodactyla domesticated         4.0         0.7
## 6 Bradypus herbi    Pilosa        <NA>        14.4         2.2
##  sleep_cycle awake brainwt  bodywt
## 1          NA  11.9      NA  50.000
## 2          NA   7.0 0.01550   0.480
## 3          NA   9.6      NA   1.350
## 4  0.1333333   9.1 0.00029   0.019
## 5  0.6666667  20.0 0.42300  600.000
## 6  0.7666667   9.6      NA   3.850
```

To select a range of columns by name, use the “:” (colon) operator

```
head(select(msleep, name:order))
```

```
##              name      genus vore      order
## 1          Cheetah  Acinonyx carni   Carnivora
## 2      Owl monkey    Aotus  omni   Primates
## 3  Mountain beaver Aplodontia herbi   Rodentia
## 4 Greater short-tailed shrew  Blarina  omni Soricomorpha
## 5              Cow      Bos herbi Artiodactyla
## 6  Three-toed sloth  Bradypus herbi    Pilosa
```

To select all columns that start with the character string “sl”, use the function `starts_with()`

```
head(select(msleep, starts_with("sl")))
```

```
##   sleep_total sleep_rem sleep_cycle
## 1         12.1        NA          NA
## 2         17.0         1.8          NA
## 3         14.4         2.4          NA
## 4         14.9         2.3  0.1333333
## 5          4.0         0.7  0.6666667
## 6         14.4         2.2  0.7666667
```

Some additional options to select columns based on a specific criteria include

1. `ends_with()` = Select columns that end with a character string
2. `contains()` = Select columns that contain a character string
3. `matches()` = Select columns that match a regular expression
4. `one_of()` = Select columns names that are from a group of names

Selecting rows using `filter()`

Filter the rows for mammals that sleep a total of more than 16 hours.

```
filter(msleep, sleep_total >= 16)
```

```
##           name      genus  vore      order conservation
## 1      Owl monkey    Aotus  omni    Primates      <NA>
## 2 Long-nosed armadillo Dasypus  carni  Cingulata      lc
## 3 North American Opossum Didelphis  omni  Didelphimorphia      lc
## 4      Big brown bat  Eptesicus  insecti  Chiroptera      lc
## 5 Thick-tailed opossum Lutreolina  carni  Didelphimorphia      lc
## 6 Little brown bat    Myotis  insecti  Chiroptera  <NA>
## 7      Giant armadillo Priodontes  insecti  Cingulata      en
## 8 Arctic ground squirrel  Sperophilus  herbi  Rodentia      lc
##   sleep_total sleep_rem sleep_cycle awake brainwt bodywt
## 1         17.0         1.8         NA   7.0 0.01550  0.480
## 2         17.4         3.1  0.3833333   6.6 0.01080  3.500
## 3         18.0         4.9  0.3333333   6.0 0.00630  1.700
## 4         19.7         3.9  0.1166667   4.3 0.00030  0.023
## 5         19.4         6.6         NA   4.6      NA  0.370
## 6         19.9         2.0  0.2000000   4.1 0.00025  0.010
## 7         18.1         6.1         NA   5.9 0.08100 60.000
## 8         16.6         NA         NA   7.4 0.00570  0.920
```

Filter the rows for mammals that sleep a total of more than 16 hours *and* have a body weight of greater than 1 kilogram.

```
filter(msleep, sleep_total >= 16, bodywt >= 1)
```

##		name	genus	vore	order	conservation
## 1	Long-nosed armadillo	Dasytus	carni	Cingulata	lc	
## 2	North American Opossum	Didelphis	omni	Didelphimorphia	lc	
## 3	Giant armadillo	Priodontes	insecti	Cingulata	en	
##	sleep_total	sleep_rem	sleep_cycle	awake	brainwt	bodywt
## 1	17.4	3.1	0.3833333	6.6	0.0108	3.5
## 2	18.0	4.9	0.3333333	6.0	0.0063	1.7
## 3	18.1	6.1	NA	5.9	0.0810	60.0

Filter the rows for mammals in the Perissodactyla and Primates taxonomic order

```
filter(msleep, order %in% c("Perissodactyla", "Primates"))
```

##		name	genus	vore	order	conservation
## 1	Owl monkey	Aotus	omni	Primates	<NA>	
## 2	Grivet	Cercopithecus	omni	Primates	lc	
## 3	Horse	Equus	herbi	Perissodactyla	domesticated	
## 4	Donkey	Equus	herbi	Perissodactyla	domesticated	
## 5	Patas monkey	Erythrocebus	omni	Primates	lc	
## 6	Galago	Galago	omni	Primates	<NA>	
## 7	Human	Homo	omni	Primates	<NA>	
## 8	Mongoose lemur	Lemur	herbi	Primates	vu	
## 9	Macaque	Macaca	omni	Primates	<NA>	
## 10	Slow loris	Nyctibeus	carni	Primates	<NA>	
## 11	Chimpanzee	Pan	omni	Primates	<NA>	
## 12	Baboon	Papio	omni	Primates	<NA>	
## 13	Potto	Perodicticus	omni	Primates	lc	
## 14	Squirrel monkey	Saimiri	omni	Primates	<NA>	
## 15	Brazilian tapir	Tapirus	herbi	Perissodactyla	vu	
##	sleep_total	sleep_rem	sleep_cycle	awake	brainwt	bodywt
## 1	17.0	1.8	NA	7.0	0.0155	0.480
## 2	10.0	0.7	NA	14.0	NA	4.750
## 3	2.9	0.6	1.0000000	21.1	0.6550	521.000
## 4	3.1	0.4	NA	20.9	0.4190	187.000
## 5	10.9	1.1	NA	13.1	0.1150	10.000
## 6	9.8	1.1	0.5500000	14.2	0.0050	0.200
## 7	8.0	1.9	1.5000000	16.0	1.3200	62.000
## 8	9.5	0.9	NA	14.5	NA	1.670
## 9	10.1	1.2	0.7500000	13.9	0.1790	6.800
## 10	11.0	NA	NA	13.0	0.0125	1.400
## 11	9.7	1.4	1.4166667	14.3	0.4400	52.200
## 12	9.4	1.0	0.6666667	14.6	0.1800	25.235
## 13	11.0	NA	NA	13.0	NA	1.100
## 14	9.6	1.4	NA	14.4	0.0200	0.743
## 15	4.4	1.0	0.9000000	19.6	0.1690	207.501

You can use the boolean operators (e.g. >, <, >=, <=, !=, %in%) to create the logical tests.

Pipe operator: %>%

Before we go any futher, let's introduce the pipe operator: %>%. dplyr imports this operator from another package (magrittr). This operator allows you to pipe the output from one function to the input of another function. Instead of nesting functions (reading from the inside to the outside), the idea of piping is to read the

functions from left to right.

Here's an example you have seen:

```
head(select(msleep, name, sleep_total))
```

```
##              name sleep_total
## 1          Cheetah      12.1
## 2        Owl monkey      17.0
## 3    Mountain beaver      14.4
## 4 Greater short-tailed shrew  14.9
## 5              Cow       4.0
## 6    Three-toed sloth      14.4
```

Now in this case, we will pipe the `msleep` data frame to the function that will select two columns (`name` and `sleep_total`) and then pipe the new data frame to the function `head()` which will return the head of the new data frame.

```
msleep %>%
  select(name, sleep_total) %>%
  head
```

```
##              name sleep_total
## 1          Cheetah      12.1
## 2        Owl monkey      17.0
## 3    Mountain beaver      14.4
## 4 Greater short-tailed shrew  14.9
## 5              Cow       4.0
## 6    Three-toed sloth      14.4
```

You will soon see how useful the pipe operator is when we start to combine many functions.

Back to dplyr verbs in action

Now that you know about the pipe operator (`%>%`), we will use it throughout the rest of this tutorial.

Arrange or re-order rows using `arrange()`

To arrange (or re-order) rows by a particular column such as the taxonomic order, list the name of the column you want to arrange the rows by

```
msleep %>% arrange(order) %>% head
```

##	name	genus	vore	order	conservation	sleep_total	sleep_rem
## 1	Tenrec	Tenrec	omni	Afrosoricida	<NA>	15.6	2.3
## 2	Cow	Bos	herbi	Artiodactyla	domesticated	4.0	0.7
## 3	Roe deer	Capreolus	herbi	Artiodactyla	lc	3.0	NA
## 4	Goat	Capri	herbi	Artiodactyla	lc	5.3	0.6
## 5	Giraffe	Giraffa	herbi	Artiodactyla	cd	1.9	0.4
## 6	Sheep	Ovis	herbi	Artiodactyla	domesticated	3.8	0.6

##	sleep_cycle	awake	brainwt	bodywt
## 1	NA	8.4	0.0026	0.900
## 2	0.6666667	20.0	0.4230	600.000
## 3	NA	21.0	0.0982	14.800
## 4	NA	18.7	0.1150	33.500
## 5	NA	22.1	NA	899.995
## 6	NA	20.2	0.1750	55.500

Now, we will select three columns from `msleep`, arrange the rows by the taxonomic order and then arrange the rows by `sleep_total`. Finally show the head of the final data frame

```
msleep %>%
  select(name, order, sleep_total) %>%
  arrange(order, sleep_total) %>%
  head
```

##	name	order	sleep_total
## 1	Tenrec	Afrosoricida	15.6
## 2	Giraffe	Artiodactyla	1.9
## 3	Roe deer	Artiodactyla	3.0
## 4	Sheep	Artiodactyla	3.8
## 5	Cow	Artiodactyla	4.0
## 6	Goat	Artiodactyla	5.3

Same as above, except here we filter the rows for mammals that sleep for 16 or more hours instead of showing the head of the final data frame

```
msleep %>%
  select(name, order, sleep_total) %>%
  arrange(order, sleep_total) %>%
  filter(sleep_total >= 16)
```

##	name	order	sleep_total
## 1	Big brown bat	Chiroptera	19.7
## 2	Little brown bat	Chiroptera	19.9
## 3	Long-nosed armadillo	Cingulata	17.4
## 4	Giant armadillo	Cingulata	18.1
## 5	North American Opossum	Didelphimorphia	18.0
## 6	Thick-tailed opossum	Didelphimorphia	19.4
## 7	Owl monkey	Primates	17.0
## 8	Arctic ground squirrel	Rodentia	16.6

Something slightly more complicated: same as above, except arrange the rows in the `sleep_total` column in a descending order. For this, use the function `desc()`

```
msleep %>%
  select(name, order, sleep_total) %>%
  arrange(order, desc(sleep_total)) %>%
  filter(sleep_total >= 16)
```

##		name	order	sleep_total
## 1		Little brown bat	Chiroptera	19.9
## 2		Big brown bat	Chiroptera	19.7
## 3		Giant armadillo	Cingulata	18.1
## 4		Long-nosed armadillo	Cingulata	17.4
## 5		Thick-tailed opossum	Didelphimorphia	19.4
## 6		North American Opossum	Didelphimorphia	18.0
## 7		Owl monkey	Primates	17.0
## 8		Arctic ground squirrel	Rodentia	16.6

Create new columns using `mutate()`

The `mutate()` function will add new columns to the data frame. Create a new column called `rem_proportion` which is the ratio of rem sleep to total amount of sleep.

```
msleep %>%
  mutate(rem_proportion = sleep_rem / sleep_total) %>%
  head
```

##		name	genus	vore	order	conservation
## 1		Cheetah	Acinonyx	carni	Carnivora	lc
## 2		Owl monkey	Aotus	omni	Primates	<NA>
## 3		Mountain beaver	Aplodontia	herbi	Rodentia	nt
## 4		Greater short-tailed shrew	Blarina	omni	Soricomorpha	lc
## 5		Cow	Bos	herbi	Artiodactyla	domesticated
## 6		Three-toed sloth	Bradypus	herbi	Pilosa	<NA>

##	sleep_total	sleep_rem	sleep_cycle	awake	brainwt	bodywt	rem_proportion
## 1	12.1	NA	NA	11.9	NA	50.000	NA
## 2	17.0	1.8	NA	7.0	0.01550	0.480	0.1058824
## 3	14.4	2.4	NA	9.6	NA	1.350	0.1666667
## 4	14.9	2.3	0.1333333	9.1	0.00029	0.019	0.1543624
## 5	4.0	0.7	0.6666667	20.0	0.42300	600.000	0.1750000
## 6	14.4	2.2	0.7666667	9.6	NA	3.850	0.1527778

You can many new columns using `mutate` (separated by commas). Here we add a second column called `bodywt_grams` which is the `bodywt` column in grams.

```
msleep %>%
  mutate(rem_proportion = sleep_rem / sleep_total,
         bodywt_grams = bodywt * 1000) %>%
  head
```



```
##           name      genus vore      order conservation
## 1      Cheetah  Acinonyx carni    Carnivora          lc
## 2      Owl monkey    Aotus  omni    Primates        <NA>
## 3  Mountain beaver Aplodontia herbi    Rodentia          nt
## 4 Greater short-tailed shrew  Blarina  omni Soricomorpha          lc
## 5              Cow        Bos  herbi Artiodactyla domesticated
## 6  Three-toed sloth  Bradypus herbi      Pilosa        <NA>
##  sleep_total sleep_rem sleep_cycle awake brainwt  bodywt rem_proportion
## 1      12.1      NA      NA  11.9      NA  50.000      NA
## 2      17.0      1.8      NA   7.0 0.01550   0.480   0.1058824
## 3      14.4      2.4      NA   9.6      NA   1.350   0.1666667
## 4      14.9      2.3  0.1333333   9.1 0.00029   0.019   0.1543624
## 5       4.0      0.7  0.6666667  20.0 0.42300  600.000   0.1750000
## 6      14.4      2.2  0.7666667   9.6      NA   3.850   0.1527778
##  bodywt_grams
## 1      50000
## 2       480
## 3      1350
## 4        19
## 5     600000
## 6      3850
```

Create summaries of the data frame using `summarise()`

The `summarise()` function will create summary statistics for a given column in the data frame such as finding the mean. For example, to compute the average number of hours of sleep, apply the `mean()` function to the column `sleep_total` and call the summary value `avg_sleep`.

```
msleep %>%
  summarise(avg_sleep = mean(sleep_total))
```

```
##   avg_sleep
## 1  10.43373
```

There are many other summary statistics you could consider such `sd()`, `min()`, `max()`, `median()`, `sum()`, `n()` (returns the length of vector), `first()` (returns first value in vector), `last()` (returns last value in vector) and `n_distinct()` (number of distinct values in vector).

```
msleep %>%
  summarise(avg_sleep = mean(sleep_total),
            min_sleep = min(sleep_total),
            max_sleep = max(sleep_total),
            total = n())
```

```
##   avg_sleep min_sleep max_sleep total
## 1  10.43373      1.9     19.9     83
```

Group operations using `group_by()`

The `group_by()` verb is an important function in dplyr. As we mentioned before it's related to concept of "split-apply-combine". We literally want to split the data frame by some variable (e.g. taxonomic order), apply a function to the individual data frames and then combine the output.

Let's do that: split the msleep data frame by the taxonomic order, then ask for the same summary statistics as above. We expect a set of summary statistics for each taxonomic order.

```
msleep %>%  
  group_by(order) %>%  
  summarise(avg_sleep = mean(sleep_total),  
            min_sleep = min(sleep_total),  
            max_sleep = max(sleep_total),  
            total = n())
```

```
## Source: local data frame [19 x 5]  
##  
##      order avg_sleep min_sleep max_sleep total  
## 1 Afrosoricida 15.600000    15.6     15.6     1  
## 2 Artiodactyla  4.516667     1.9      9.1     6  
## 3 Carnivora    10.116667     3.5    15.8    12  
## 4 Cetacea      4.500000     2.7     5.6     3  
## 5 Chiroptera   19.800000    19.7    19.9     2  
## 6 Cingulata    17.750000    17.4    18.1     2  
## 7 Didelphimorphia 18.700000    18.0    19.4     2  
## 8 Diprotodontia 12.400000    11.1    13.7     2  
## 9 Erinaceomorpha 10.200000    10.1    10.3     2  
## 10 Hyracoidea   5.666667     5.3     6.3     3  
## 11 Lagomorpha   8.400000     8.4     8.4     1  
## 12 Monotremata  8.600000     8.6     8.6     1  
## 13 Perissodactyla 3.466667     2.9     4.4     3  
## 14 Pilosa      14.400000    14.4    14.4     1  
## 15 Primates    10.500000     8.0    17.0    12  
## 16 Proboscidea  3.600000     3.3     3.9     2  
## 17 Rodentia    12.468182     7.0    16.6    22  
## 18 Scandentia   8.900000     8.9     8.9     1  
## 19 Soricomorpha 11.100000     8.4    14.9     5
```